AWS Global Infrastructure

AWS Datacenters

50k to 80k Servers and required components (e.g. storage/network)

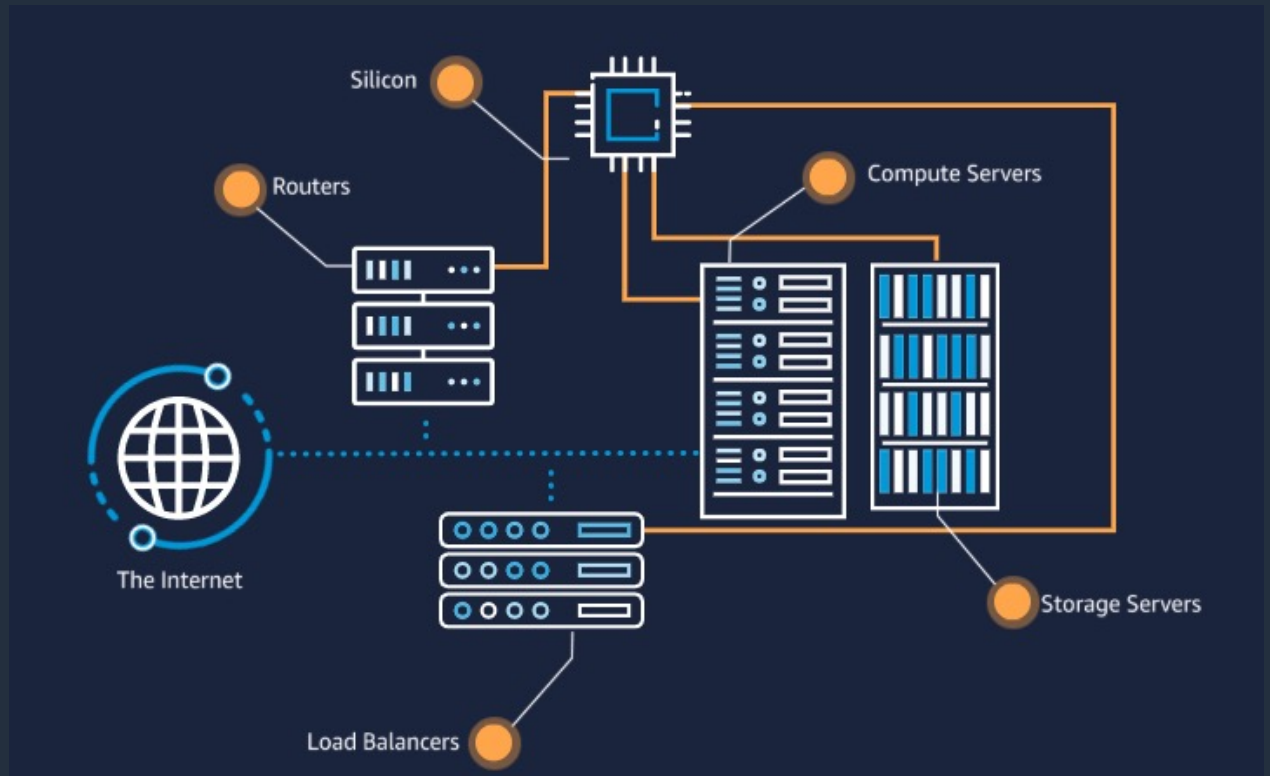

AWS Availability Zone (AZ)

One or More than one Datacenters

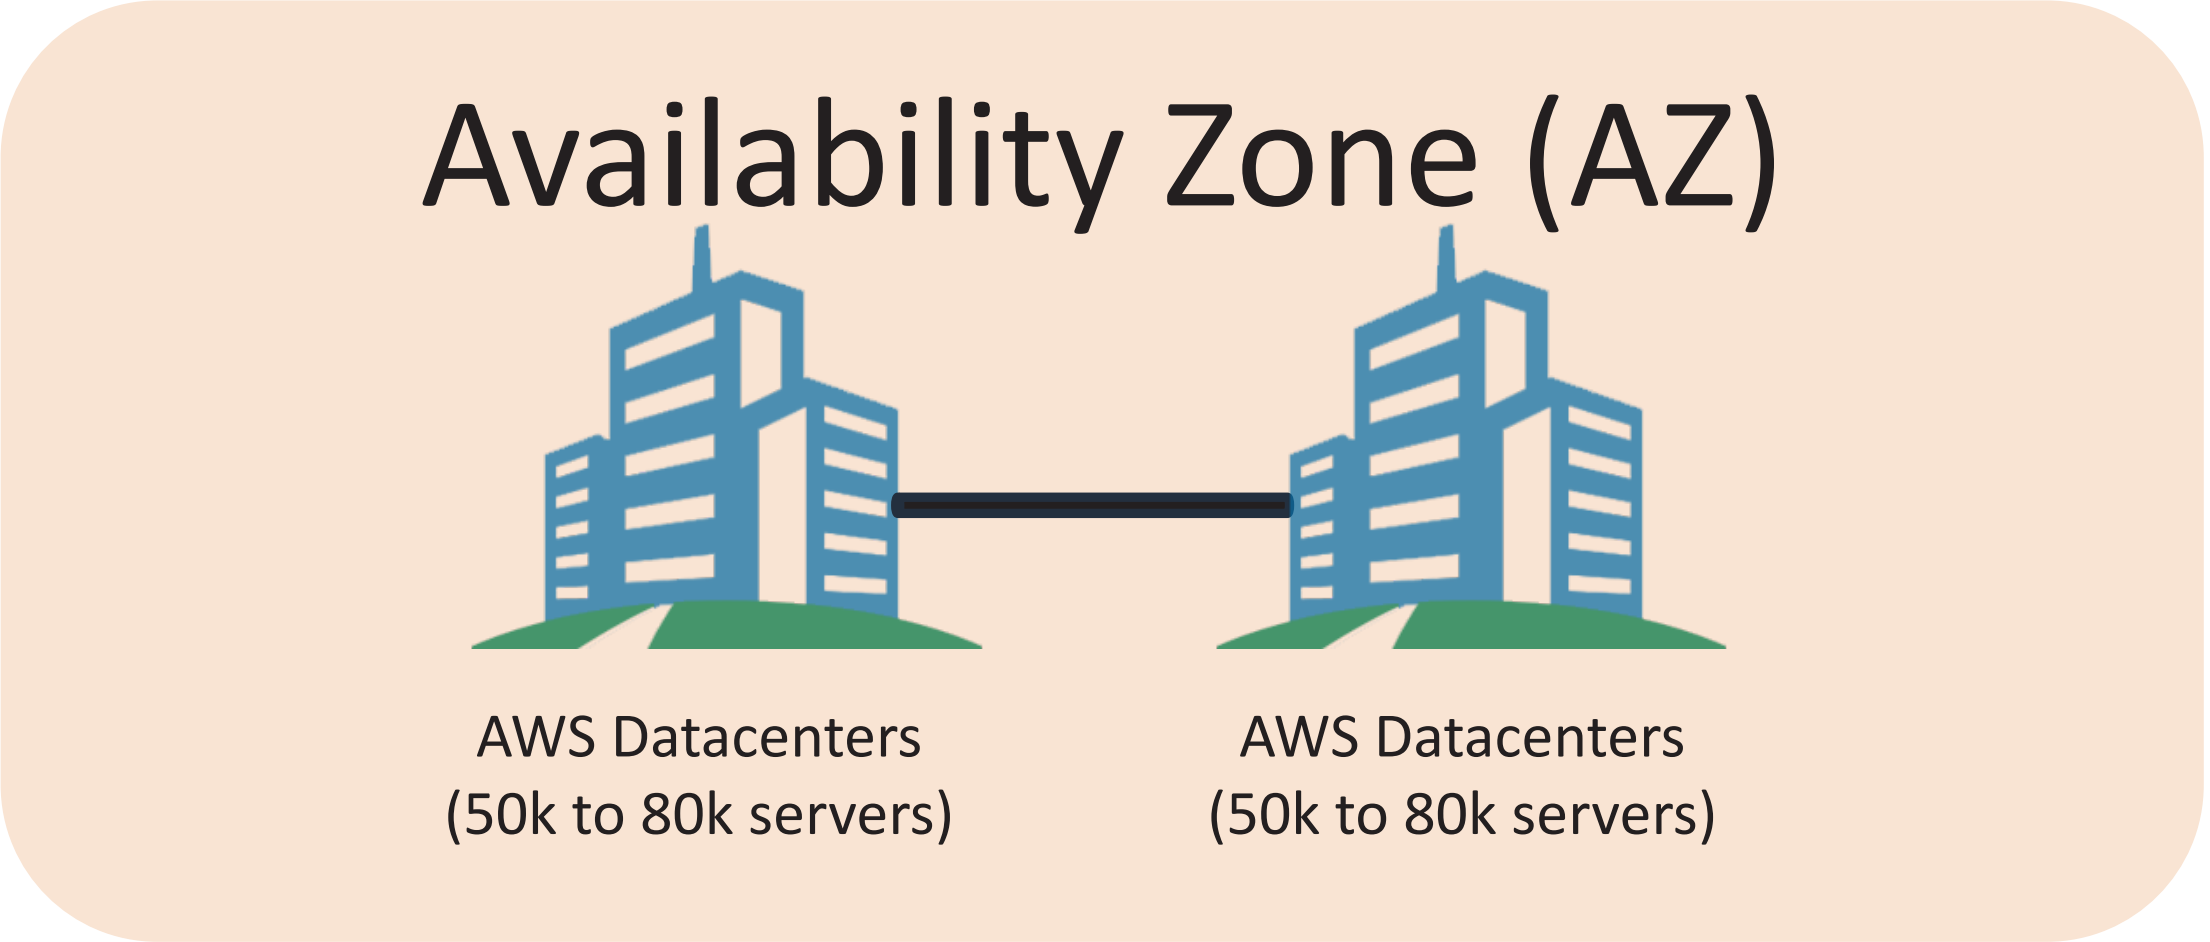

AWS Availability Zone (AZ) Design

- Fully isolated infrastructure with one or more datacenters
- Meaningful distance of separation
- Unique power infrastructure
- Many 100Ks of servers at scale
- Datacenters connected via fully redundant and isolated metro fiber

AWS Network Design

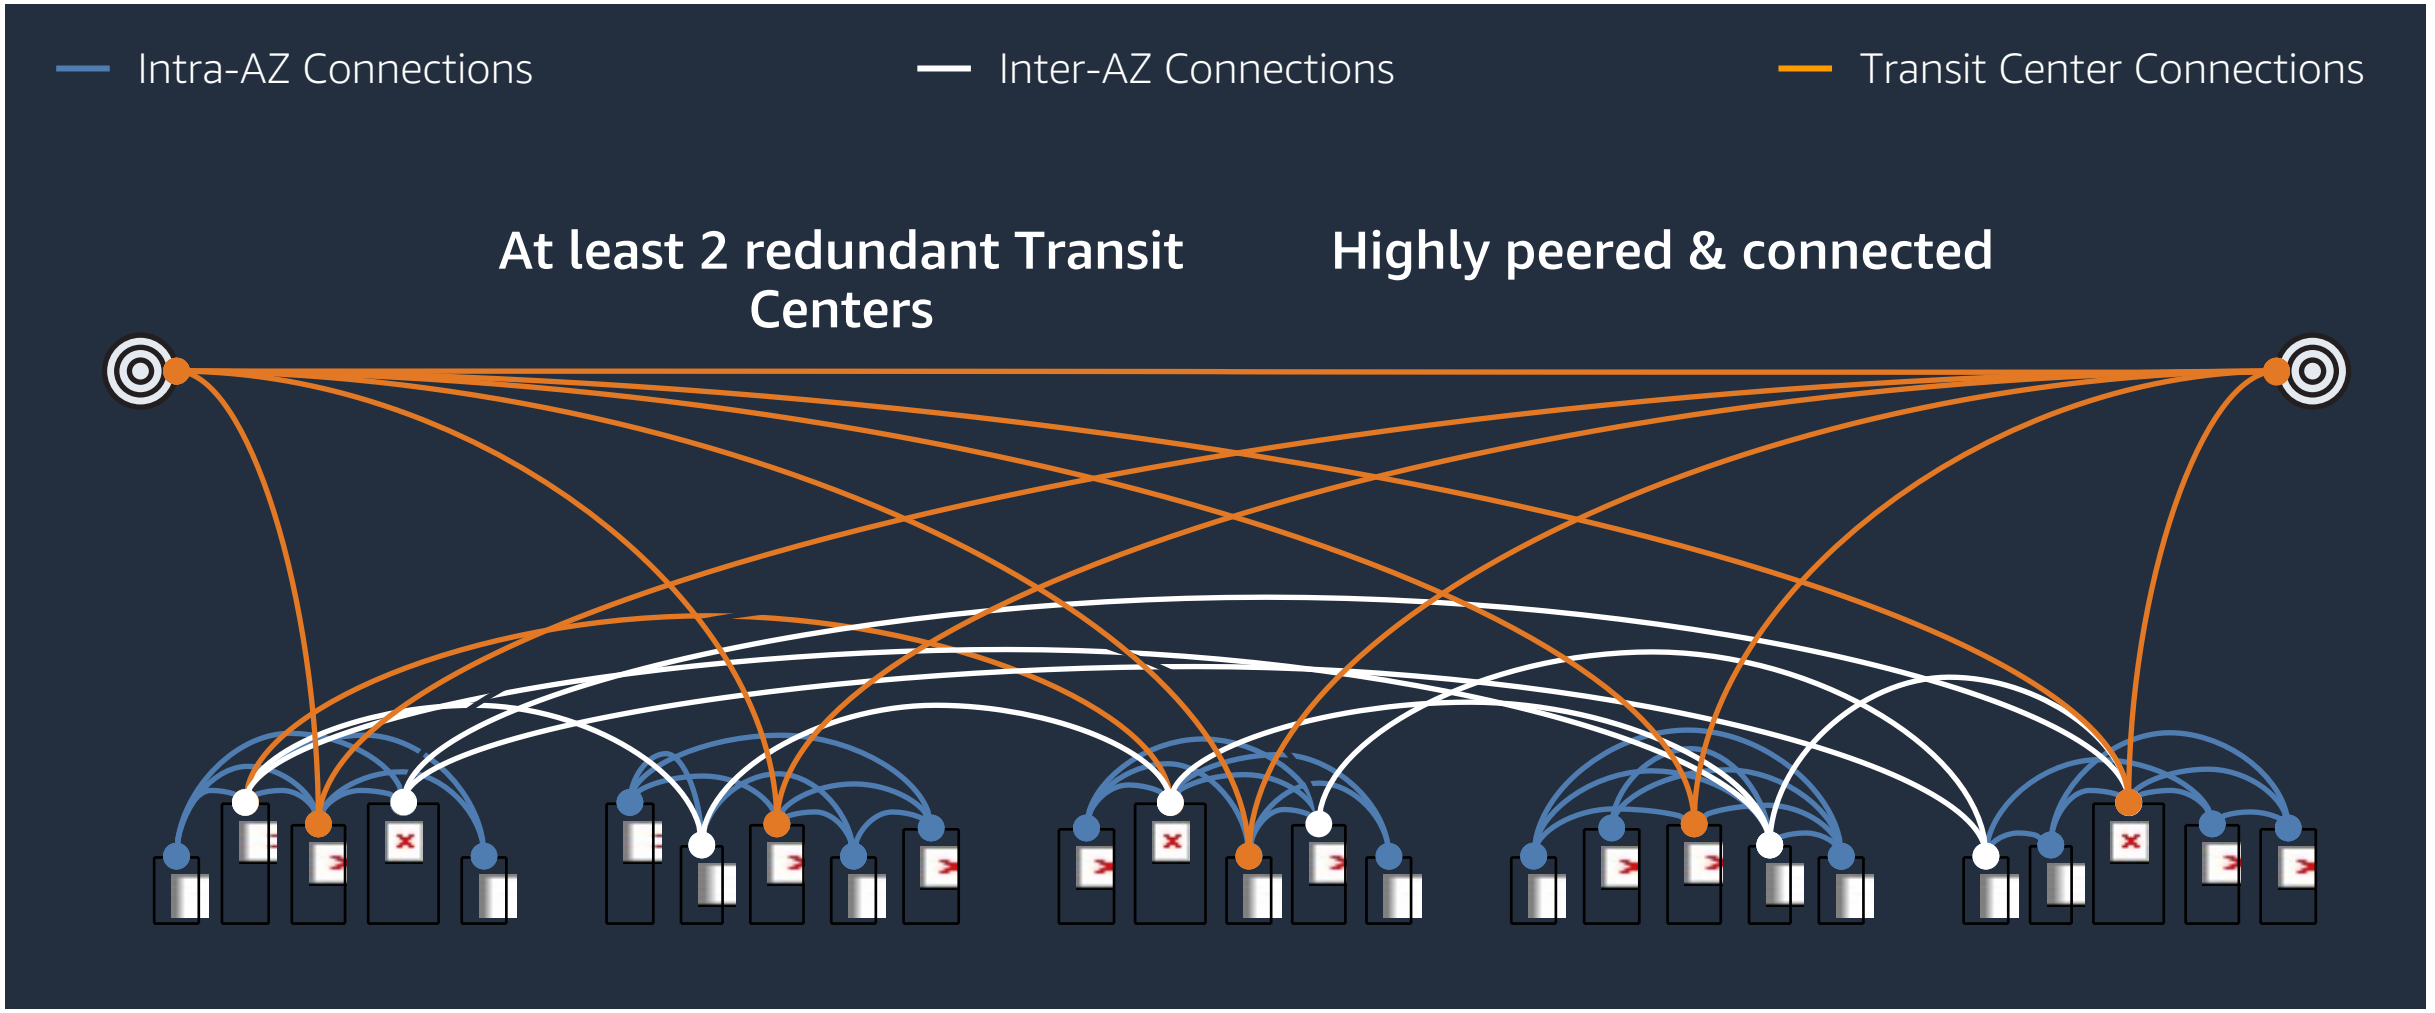

AWS Region

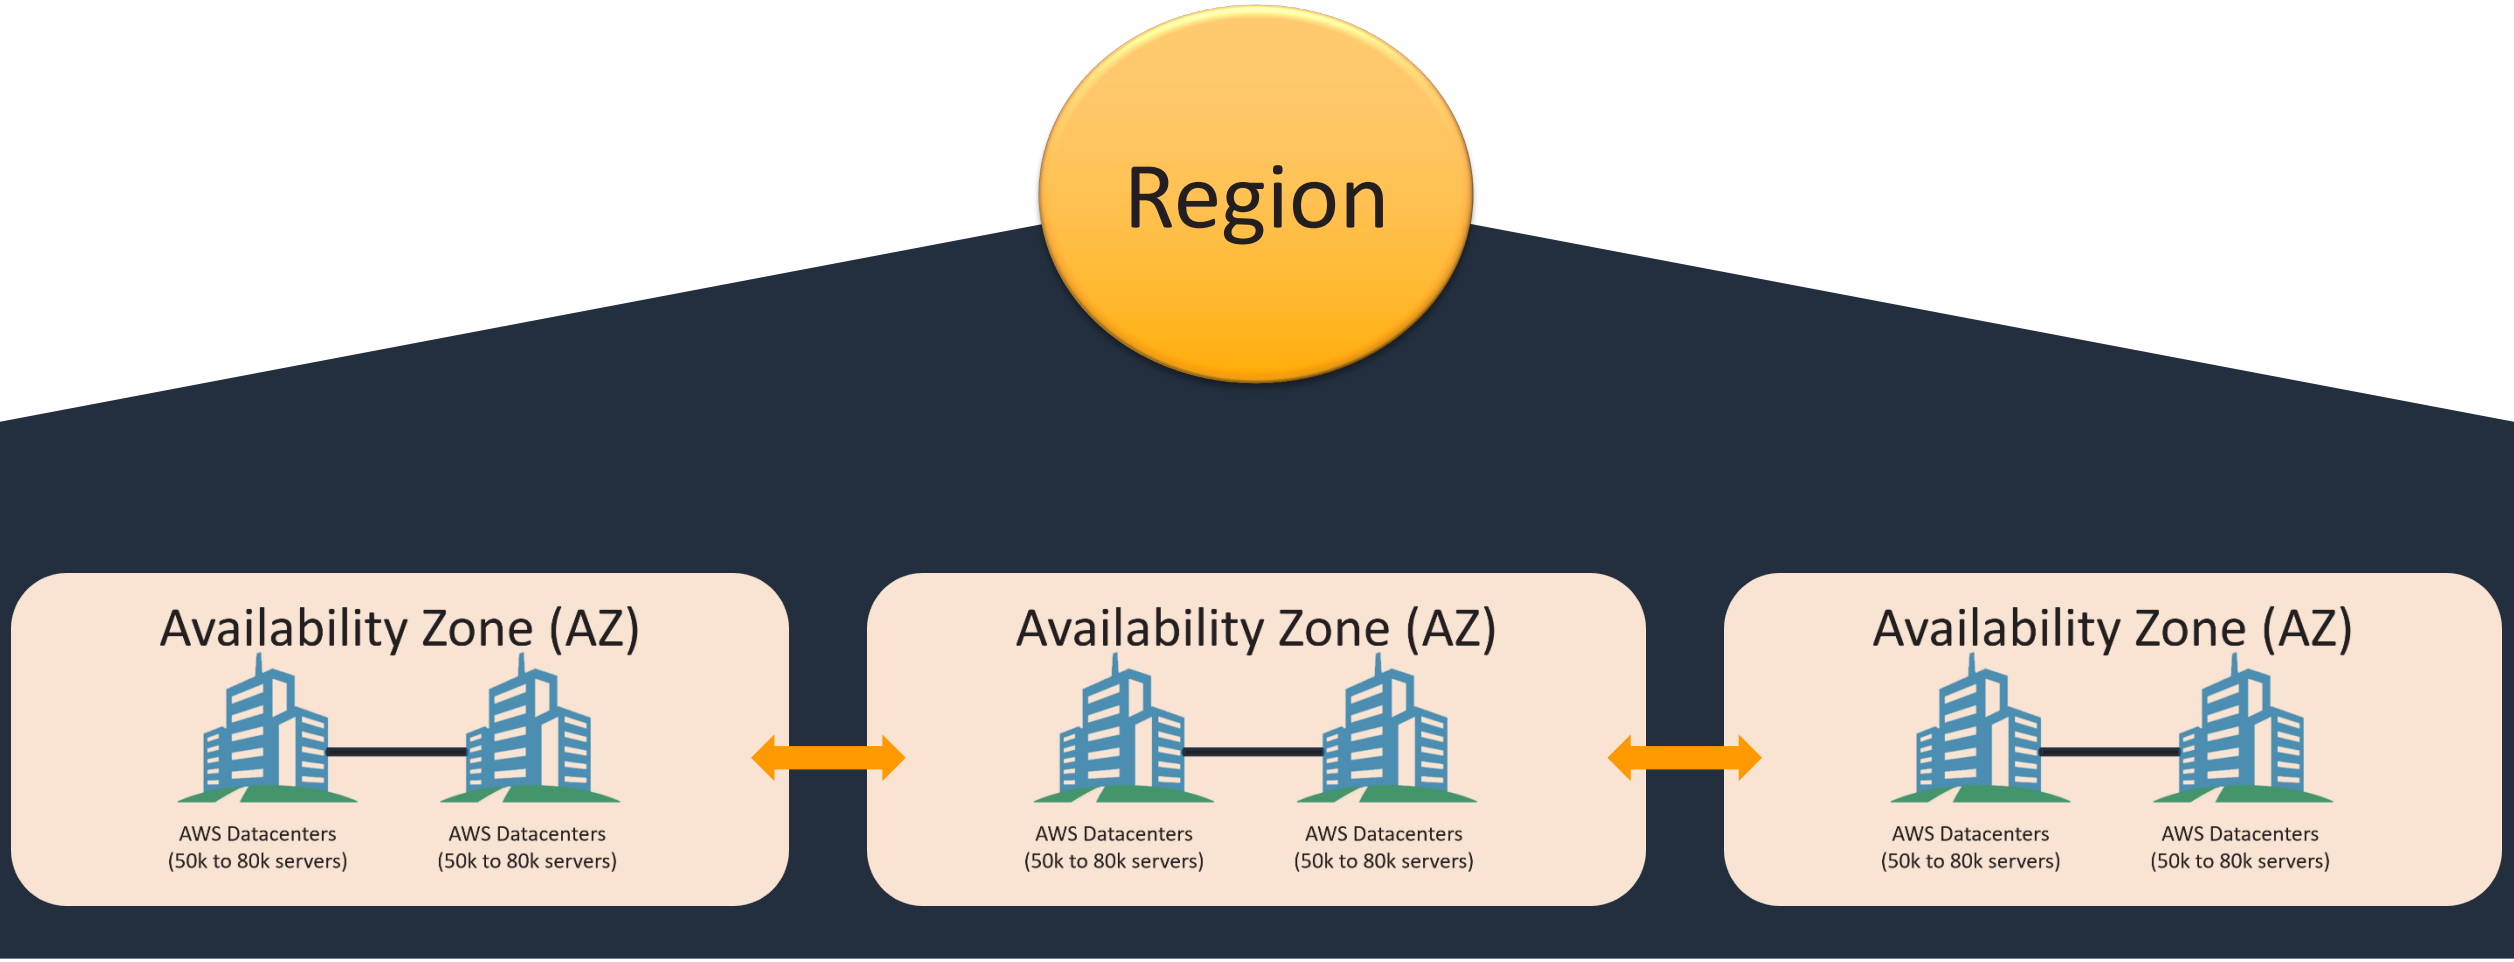

AWS Global Infrastructure

24 Launched Regions

Each with multiple Availability Zones (AZ's)

3 Announced Regions

76 Availability Zones

1 Local Zone

For ultralow latency applications

2x More Regions

With multiple AZ's than the next largest cloud provider

245 Countries and Territories Served

97 Direct Connect Locations

216 Points of Presence

205 Edge Locations and 11 Regional Edge Caches

AWS Regions

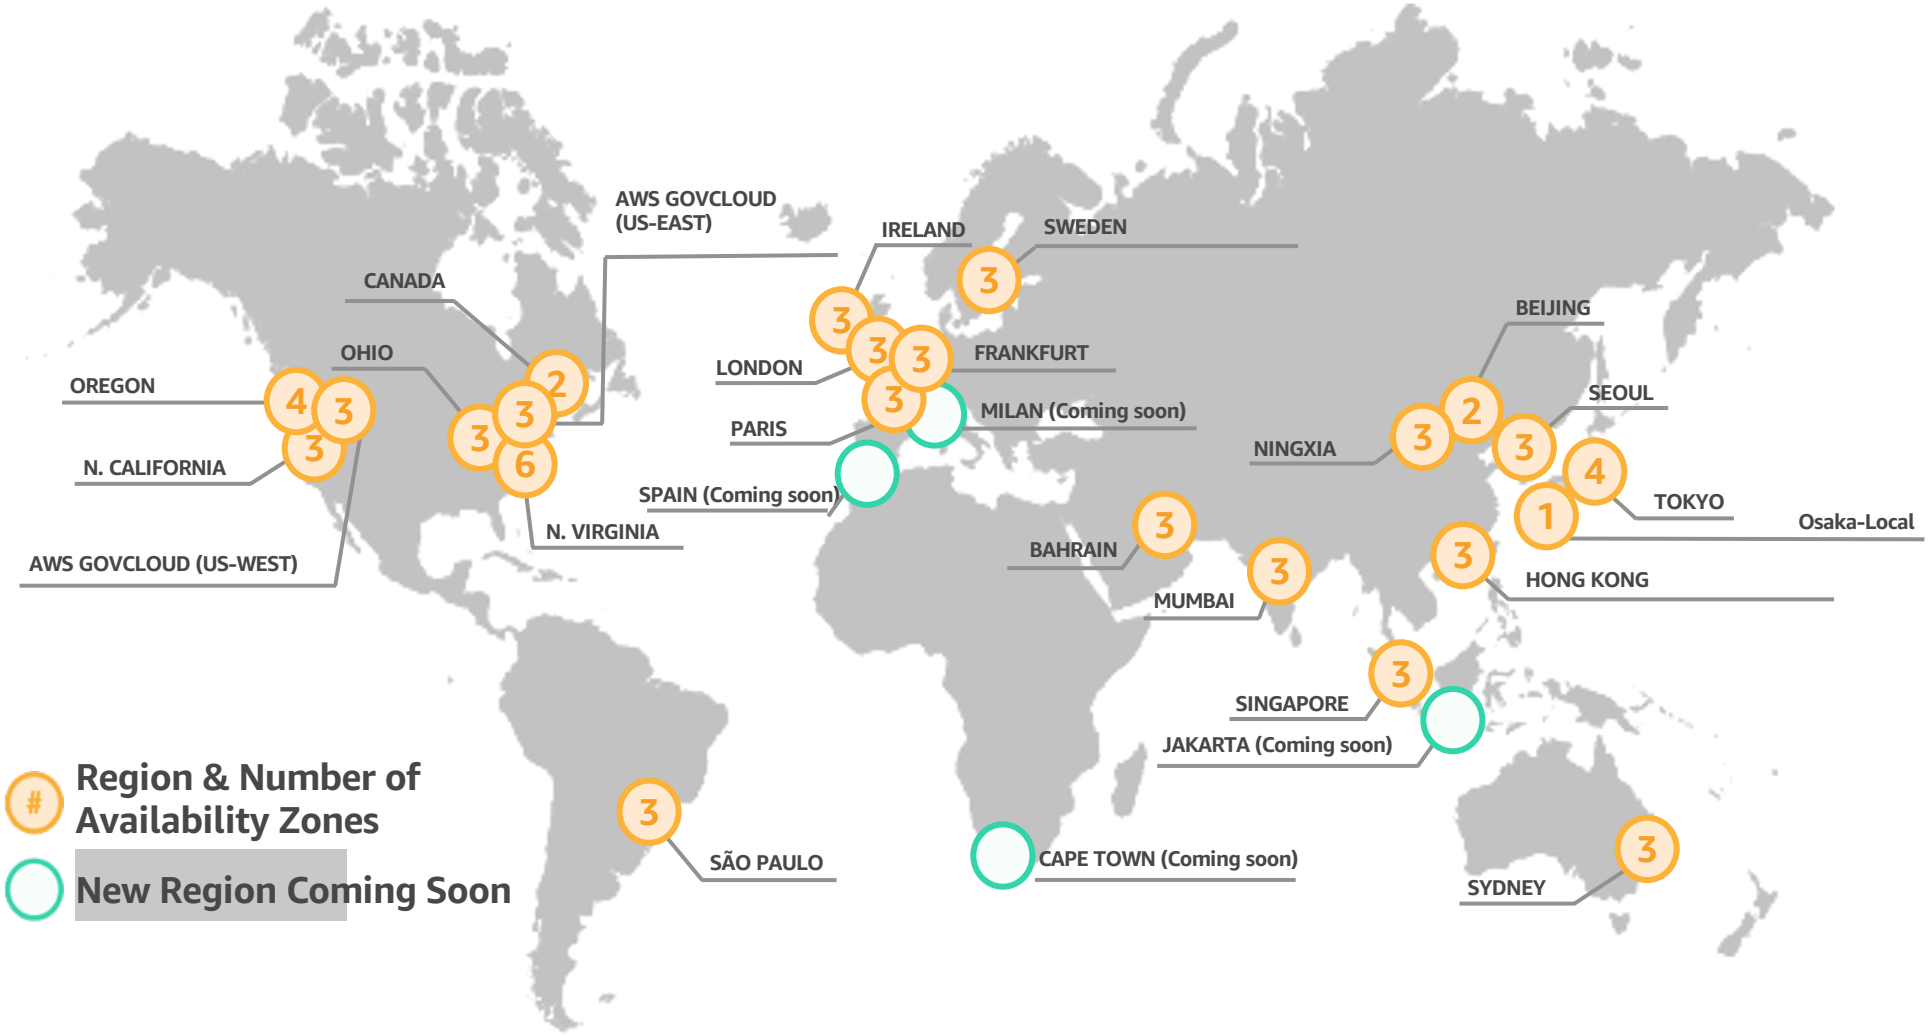

Edge Location (Point or Presence)

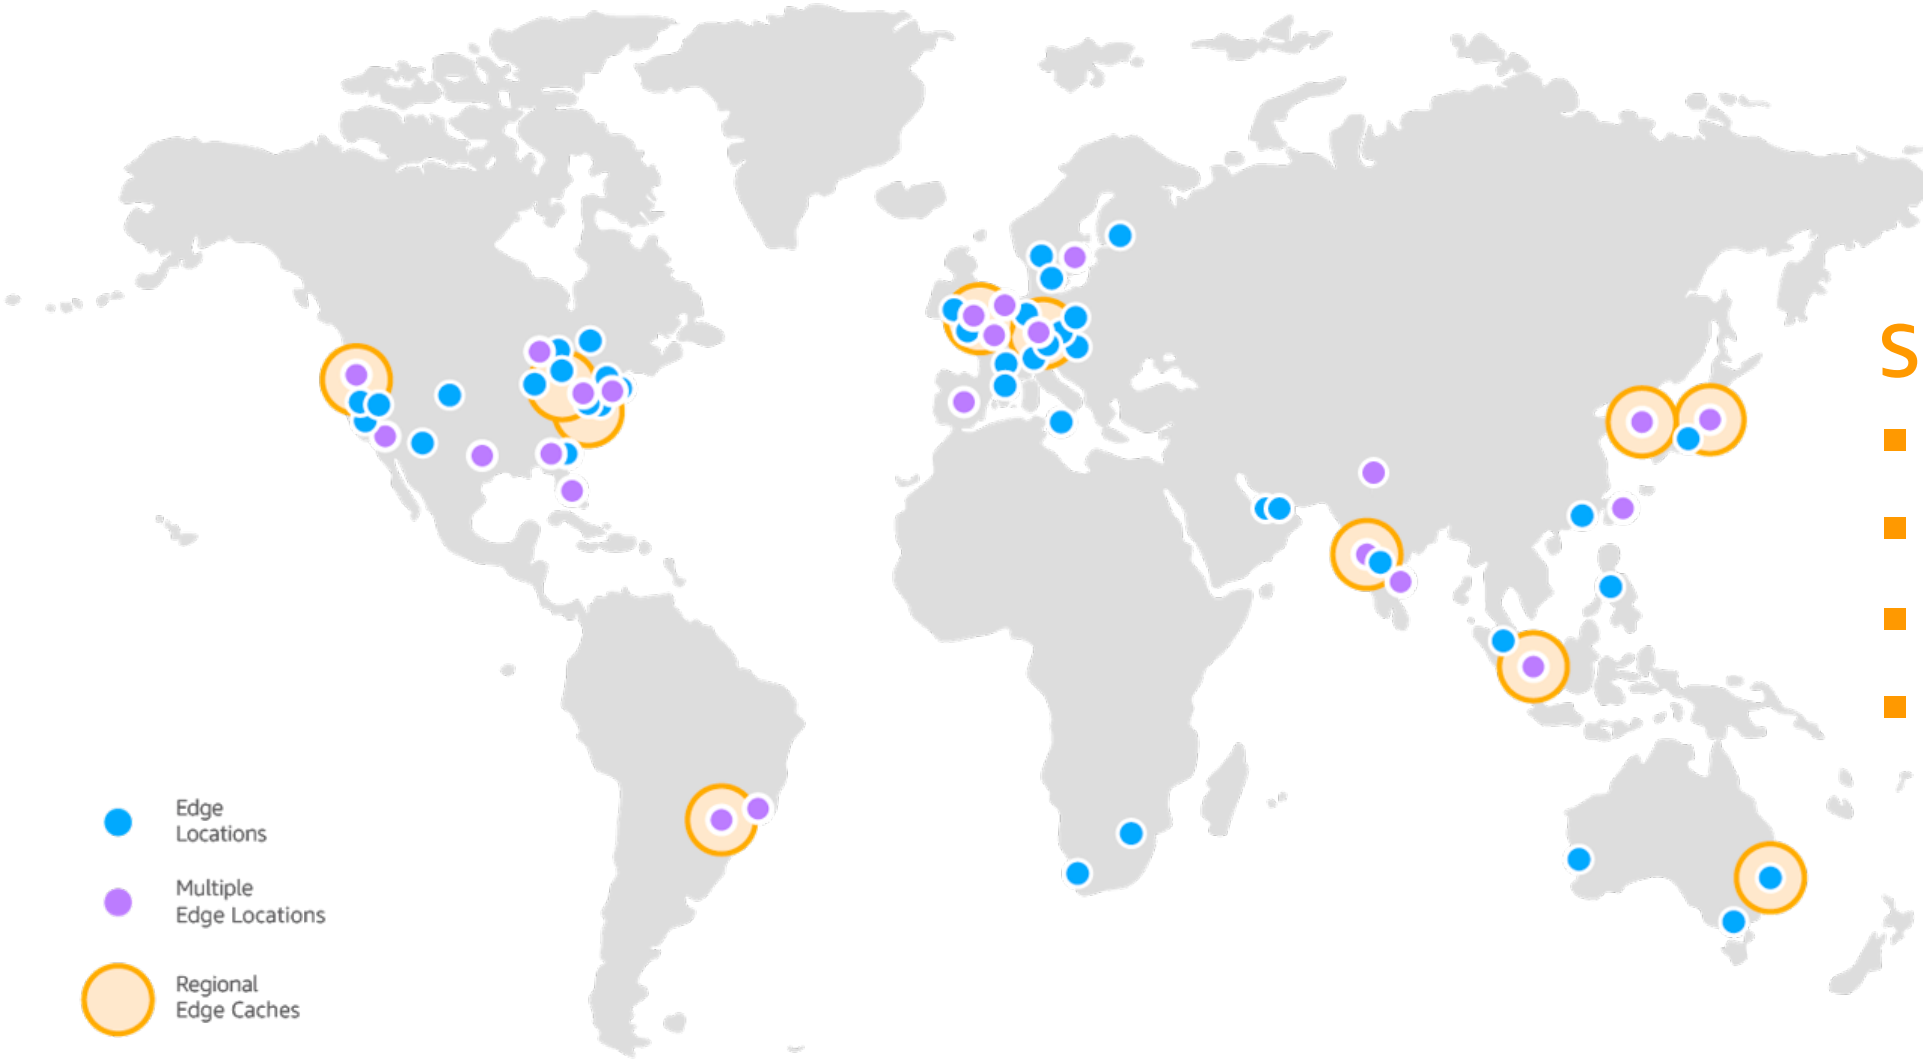

Services

- Route 53
- CloudFront
- WAF/Shield
- Lambda@Edge

aws training and certification

Choosing AWS Regions for your Architectures

Choosing a Region

aws training and certification

Data residency and regulatory compliance

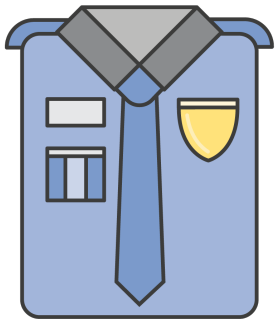

Are there relevant data privacy laws in the Region?

Can customer data be stored outside the country?

Can you meet your governance obligation?

Choosing a Region

aws training and certification

Proximity of users to data

Small differences in latency can impact customer experience

Choose the Region closest to your users

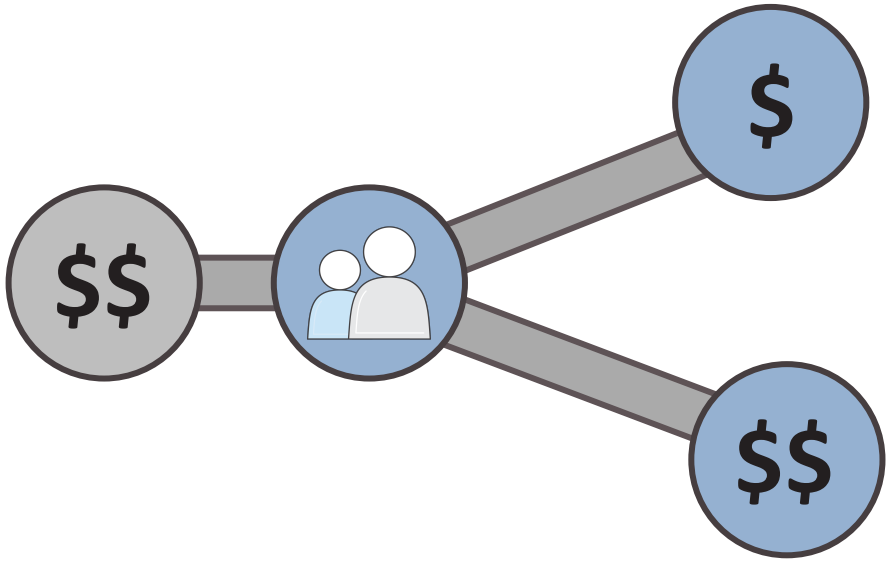

Choosing a Region

aws training and certification

Service and feature availability

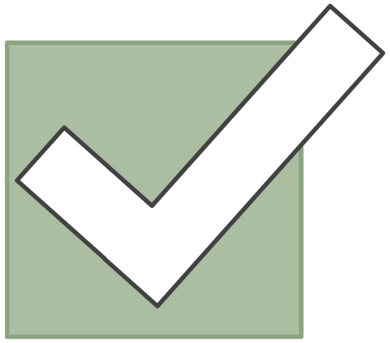

Some services not yet available in all AWS Regions

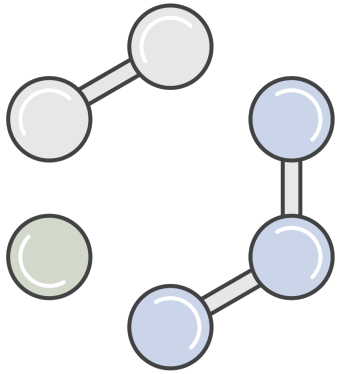

Can use some services cross-Region, but at increased latency

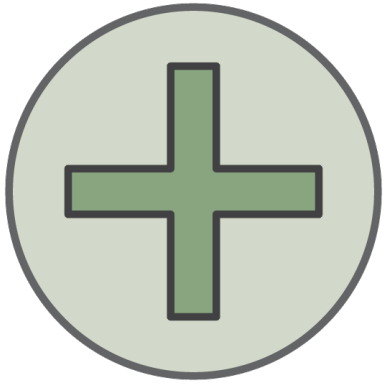

Services expanded to new Regions regularly

Choosing a Region

aws training and certification

Cost-effectiveness

- Costs vary by AWS Region
- Some services like Amazon S3 have costs for transferring data out
- Consider the cost-effectiveness of replicating the entire environment in another Region